

# INTELLIGENT CANDIDATE DISCOVERY & RANKING

A Hybrid Rules-Engine and Semantic Ranking Pipeline for Founding AI Engineers

Prepared by: @dvenkatareddy2006\_4138  
Redrob Data & AI Challenge Solution

# The Core Challenge: Beyond Keywords

---

## THE KEYWORD-STUFFER TRAP

- Traditional keyword search is easily gamed by profiles that dump buzzwords into a skill list without actual experience.
- Example: A 'Marketing Manager' listing 20 AI keywords (like LLMs, RAG, Pinecone) is auto-ranked highly by naive parsers.
- Standard LLM-per-candidate evaluation pipelines are cost-prohibitive and fail compute latency restrictions on 100K pools.

## OUR MULTI-LAYER APPROACH

- Title-as-a-Gate: Treats current and past titles as ceiling multipliers. Irrelevant roles are capped immediately.
- Profile Data Consistency: Validates internal alignment (years of experience vs. total employment span duration).
- Reachability Integration: Incorporates platform response rates, activity dates, and notice periods as bounded multipliers.

# System Architecture: Streamlined Pipeline

---

## 1. Stream & Filter

[features.py]

Reads candidates.jsonl line-by-line. Identifies and isolates honeypots/fake accounts instantly.

## 2. Feature Extract

[features.py]

Extracts flat feature vectors, computes soft disqualifiers (consulting-only, pure research).

## 3. Hybrid Score

[scoring.py]

Combines title gate, exact skill matching, semantic similarity, and behavioral multipliers.

## 4. Reason & Rank

[reasoning.py]

Generates human-auditable reasons. Deterministically sorts and enforces monotonic ranks.

# Trap Defense: Honeytrap & Disqualifiers

---

## HARD HONEYPOT EXCLUSIONS

- Zero-Duration Experts: Excludes candidates claiming 'Expert' level in 3+ skills but showing 0 months of use.
- Stated Experience Mismatch: Excludes profiles where the sum of employment durations exceeds the stated total experience by >15%.
- Overlapping Jobs: Flags and removes profiles claiming multiple concurrent full-time jobs (transition window padding excluded).

## SOFT DISQUALIFIERS

- Consulting Only: Profiles where 100% of career history lies in outsourcing/IT-consulting giants, which the JD explicitly flags.
- Pure Academic/Research: Profiles with no industry product/deployment background, focusing purely on research lab/academia.
- Title Chasing: Rapid career escalation (e.g. Lead/Staff titles) with sub-18-month tenures across multiple companies.

# The Hybrid Scoring Engine

---

## SCORING BREAKDOWN

- Title Fit: 22% - Core vs Adjacent titles
- Skill Match: 23% - Mandatory group coverage
- Career Evidence: 20% - Shipped-system references
- Experience Band: 10% - Sweet spot (5-9 Years)
- Location/Edu: 10% - Tier-1 cities / Tier-1 edu
- Redrob Behavior: 15% - Platform engagement multiplier

## CORE ALGORITHMIC FEATURES

- Dynamic Normalization: If a candidate has missing fields (e.g. no educational tier or no github\_activity\_score), weights automatically redistribute proportionally so candidates are not zero-penalized.
- Semantic Embeddings: Blends exact token matching with cosine similarity from a local sentence-transformer (all-MiniLM-L6-v2) to match synonyms (e.g. 'FAISS' or 'BM25' to indexing concepts).
- Behavioral Multiplier: Acts as a multiplier with a 0.55x floor to prevent low responsiveness from tanking a strong profile, but prioritizing active searchers.

# Audit Trails & Reasoning Column

---

## HUMAN-READABLE DECISION JUSTIFICATION

- **Deterministic Template Composer:** Uses real profile facts (years of experience, title, current employer, specific matching skills) to compose the reasoning column, ensuring zero hallucinations.
- **Linguistic Variation:** Uses candidate\_id hash-seeding to choose between multiple equivalent grammatical templates. This creates natural-reading text across rows, satisfying the hackathon's verification audits.
- **Clear Conflict Signaling:** If the ranker down-ranks a candidate due to soft disqualifiers (e.g. a long notice period or a consulting-firm background), the concern is explicitly outputted in the justification for human audit.
- **Example output:** 'Senior AI Engineer at TechCorp (7.5 yrs) with hands-on embeddings retrieval, vector db background. Lists PyTorch, Pinecone directly. Concern: entire career at consulting firms.'

# Compute Compliance & Performance Matrix

## COMPUTE SUMMARY

- Strict Compliance: Meets all restrictions listed in submission\_spec.docx.
- Streaming Parser: Avoids full in-memory loads. Processes 100K rows in O(1) memory space.
- Fully Reproducible: Executed and verified inside a CPU-only sandbox, generating the top 100 CSV in under 40s.

| Metric  | Required Limit     | Our Result   |
|---------|--------------------|--------------|
| Runtime | ≤ 5 mins           | 38.1 seconds |
| Memory  | ≤ 16 GB            | < 1 GB       |
| Network | Disabled (Offline) | 100% Offline |
| GPU     | None               | CPU Only     |